


(a) Hidden-intensity recovery vs capacity

dictionary size d_{sae}

BatchTopK-SAE (per-token)

T-SAE (per-token)

TXC-pre (T=2)

TXC-pre (T=4)

TXC-pre (T=8)


TXC-post (T=2)

TXC-post (T=4)

TXC-post (T=8)

Stacked-SAE (T=2)

(b) Recovery rises with T, saturates by T = 4



BatchTopK-SAE (per-token)

T-SAE (per-token)

TXC-pre (window)

TXC-post (window)

Stacked-SAE (window)

window size T (T = 1: per-token)

Stacked-SAE (T=4)

Stacked-SAE (T=8)