G1 PORT-FIDELITY: NO FORMAL VERDICT POSTED — the original pilot is VOID as a port test (wrong stream)
substrate confirmed base-l12; warmup=0 closed the schedule gap (descent ratio 3.63→1.30 vs upstream)

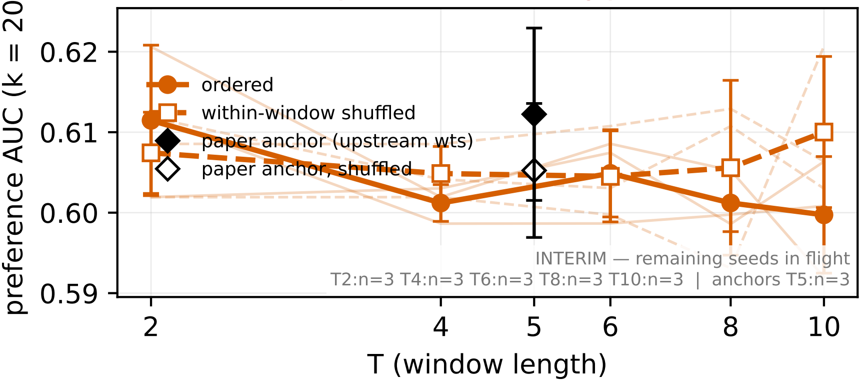

paper-faithful arm: agentic_txc_02_v1t = the paper's own RLHF TXC architecture (MatryoshkaTXCDRContrastiveMultiscale) ported to this harness, on the paper stream (gemma_2_2b_base_l12_phase7). Conclusions here are statements about the PAPER arm itself, not the plain-TXC modernization (that is the companion figure). Port-fidelity status is stamped above the axes — read it before quoting these numbers as the paper's. KNOWN DEVIATIONS from the upstream recipe, disclosed not patched: (i) upstream clips gradients at 1.0; this harness applies NO gradient clipping. (ii) rows record precision='bf16' but training is fp32 throughout — the field is declarative only (no autocast/GradScaler). Everything else — const Adam 3e-4, no warmup, batch schedule, k_win=100·T, step budget — follows upstream.