


Ordering declining and helping inside one response: the useful control is a schedule, not a level
Qwen-2.5-1.5B, layer 14 · target and foil hold the SAME sentences in a different order · bars = SEM, n = 32 pairs (A), 160 slots (B)

Δ log-prob margin toward the intended stance order
(steered – unsteered, summed over the response)

A. Steering the refuse/comply order within one response
(target and foil are the same sentences, reordered)



fraction of sentence slots whose declining-vs-helping preference moved the intended way

B. Behavioural check: does the model's own choice follow?
(held-out candidates, model's intrinsic preference differenced out)

