


Does a temporal crosscoder steer trajectories better than a TopK SAE? Not for temporal reasons.
Qwen-2.5-1.5B layer 14 · both dictionaries trained on one activation cache at matched token-activations/step · d_sae 4096 · n=20 pairs

A. Destroying the learned temporal order costs nothing
(learned content matters; its arrangement in time does not)



B. Counted in scalars rather than latents,
the SAE is not behind

